

Manual pages

btrfs(8)

btrfs(5)

btrfs-balance(8)

btrfs-check(8)

btrfs-convert(8)

btrfs-device(8)

SYNOPSIS

DESCRIPTION

DEVICE MANAGEMENT

Typical use cases

SUBCOMMAND

DEVICE STATS

EXIT STATUS

AVAILABILITY

SEE ALSO

btrfs-filessystem(8)

btrfs-find-root(8)

btrfs-image(8)

btrfs-inspect-internal(8)

btrfs-ioctl(2)

btrfs-map-logical(8)

btrfs-property(8)

btrfs-qgroup(8)

btrfs-quota(8)

btrfs-receive(8)

btrfs-replace(8)

btrfs-rescue(8)

btrfs-restore(8)

btrfs-scrub(8)

btrfs-select-super(8)

btrfs-send(8)

btrfs-subvolume(8)

btrfstune(8)

fsck.btrfs(8)

mkfs.btrfs(8)

Administration

Hardware considerations

Changes (feature/version)

Changes (kernel/version)

Changes (btrfs-progs)

Contributors

Glossary

Installation instructions

Source repositories

Interoperability

FEATURES

Common Linux features

Custom ioctls

Auto-repair on read

Balance

Compression

Checksumming

Convert

Deduplication

Defragmentation

Inline files

Quota groups

Reflink

Resize

Scrub

Seeding device

Send/receive

Subpage support

Subvolumes

Swapfile

Tree checker

Trim/discard

Volume management

Zoned mode

DEVELOPER DOCUMENTATION

Development notes

Developer's FAQ

Conventions and style for documentation

Experimental features

Btrfs design

Btrees

On-disk Format

Send stream format

JSON output

Internal APIs

Release checklist

Pull request review workflow

Command line, formatting, UI guidelines

btrfs-ioctl(2)

TODO

Troubleshooting pages

# btrfs-device(8)

## SYNOPSIS

**btrfs device** <subcommand> <args>

## DESCRIPTION

The **btrfs device** command group is used to manage devices of the btrfs filesystems.

## DEVICE MANAGEMENT

BTRFS filesystem can be created on top of single or multiple block devices. Devices can be then added, removed or replaced on demand. Data and metadata are organized in allocation profiles with various redundancy policies. There's some similarity with traditional RAID levels, but this could be confusing to users familiar with the traditional meaning. Due to the similarity, the RAID terminology is widely used in the documentation. See [mkfs.btrfs\(8\)](#) for more details and the exact profile capabilities and constraints.

The device management works on a mounted filesystem. Devices can be added, removed or replaced, by commands provided by **btrfs device** and **btrfs replace**.

The profiles can be also changed, provided there's enough workspace to do the conversion, using the **btrfs balance** command and namely the filter *convert*.

### Type

The block group profile type is the main distinction of the information stored on the block device. User data are called *Data*, the internal data structures managed by filesystem are *Metadata* and *System*.

### Profile

A profile describes an allocation policy based on the redundancy/replication constraints in connection with the number of devices. The profile applies to data and metadata block groups separately. E.g. *single*, *RAID1*.

### RAID level

Where applicable, the level refers to a profile that matches constraints of the standard RAID levels. At the moment the supported ones are: RAID0, RAID1, RAID10, RAID5 and RAID6.

## Typical use cases

### Starting with a single-device filesystem

Assume we've created a filesystem on a block device `/dev/sda` with profile *single/single* (data/metadata), the device size is 50GiB and we've used the whole device for the filesystem. The mount point is `/mnt`.

The amount of data stored is 16GiB, metadata have allocated 2GiB.

### Add new device

We want to increase the total size of the filesystem and keep the profiles. The size of the new device `/dev/sdb` is 100GiB.

```
$ btrfs device add /dev/sdb /mnt
```

The amount of free data space increases by less than 100GiB, some space is allocated for metadata.

### Convert to RAID1

Now we want to increase the redundancy level of both data and metadata, but we'll do that in steps. Note, that the device sizes are not equal and we'll use that to show the capabilities of split data/metadata and independent profiles.

The constraint for RAID1 gives us at most 50GiB of usable space and exactly 2 copies will be stored on the devices.

First we'll convert the metadata. As the metadata occupy less than 50GiB and there's enough workspace for the conversion process, we can do:

```
$ btrfs balance start -mconvert=raid1 /mnt
```

This operation can take a while, because all metadata have to be moved and all block pointers updated. Depending on the physical locations of the old and new blocks, the disk seeking is the key factor affecting performance.

You'll note that the system block group has been also converted to RAID1, this normally happens as the system block group also holds metadata (the physical to logical mappings).

What changed:

- available data space decreased by 3GiB, usable roughly  $(50 - 3) + (100 - 3) = 144$  GiB
- metadata redundancy increased

IOW, the unequal device sizes allow for combined space for data yet improved redundancy for metadata. If we decide to increase redundancy of data as well, we're going to lose 50GiB of the second device for obvious reasons.

```
$ btrfs balance start -dconvert=raid1 /mnt
```

The balance process needs some workspace (i.e. a free device space without any data or metadata block groups) so the command could fail if there's too much data or the block groups occupy the whole first device.

The device size of `/dev/sdb` as seen by the filesystem remains unchanged, but the logical space from 50-100GiB will be unused.

### Remove device

Device removal must satisfy the profile constraints, otherwise the command fails. For example:

```
$ btrfs device remove /dev/sda /mnt
ERROR: error removing device '/dev/sda': unable to go below two devices on raid1
```

In order to remove a device, you need to convert the profile in this case:

```
$ btrfs balance start -mconvert=dup -dconvert=single /mnt
$ btrfs device remove /dev/sda /mnt
```

ⓘ Warning

Do not run balance to convert from a profile with more redundancy to one with less redundancy in order to remove a failing device from a filesystem.

Balance is done by reading out the good metadata/data and write them into into a new chunk. Thus it's possible the new chunk is written into the failing device.

Use **btrfs device replace** instead.

## SUBCOMMAND



- *Data, single, Metadata, single, System, single* -- in general, list of block group type (Data, Metadata, System) and profile (single, RAID1, ...) allocated on the device
- *Data, RAID0/3* -- in particular, striped profiles RAID0/RAID10/RAID5/RAID6 with the number of devices on which the stripes are allocated, multiple occurrences of the same profile can appear in case a new device has been added and all new available stripes have been used for writes
- *Unallocated* -- remaining space that the filesystem can still use for new block groups

Options

-b|--raw

raw numbers in bytes, without the B suffix

-h|--human-readable

print human friendly numbers, base 1024, this is the default

-H

print human friendly numbers, base 1000

--iec

select the 1024 base for the following options, according to the IEC standard

--si

select the 1000 base for the following options, according to the SI standard

-k|--kbytes

show sizes in KiB, or kB with --si

-m|--mbytes

show sizes in MiB, or MB with --si

-g|--gbytes

show sizes in GiB, or GB with --si

-t|--tbytes

show sizes in TiB, or TB with --si

If conflicting options are passed, the last one takes precedence.

## DEVICE STATS

The device stats keep persistent record of several error classes related to doing IO. The current values are printed at mount time and updated during filesystem lifetime or from a scrub run.

```
$ btrfs device stats /dev/sda3
[/dev/sda3].write_io_errs    0
[/dev/sda3].read_io_errs    0
[/dev/sda3].flush_io_errs   0
[/dev/sda3].corruption_errs 0
[/dev/sda3].generation_errs 0
```

write\_io\_errs

Failed writes to the block devices, means that the layers beneath the filesystem were not able to satisfy the write request.

read\_io\_errors

Read request analogy to write\_io\_errs.

flush\_io\_errs

Number of failed writes with the *FLUSH* flag set. The flushing is a method of forcing a particular order between write requests and is crucial for implementing crash consistency. In case of btrfs, all the metadata blocks must be permanently stored on the block device before the superblock is written.

corruption\_errs

A block checksum mismatched or a corrupted metadata header was found.

generation\_errs

The block generation does not match the expected value (e.g. stored in the parent node).

Since kernel 5.14 the device stats are also available in textual form in `/sys/fs/btrfs/FSID/devinfo/DEVID/error_stats`.

## EXIT STATUS

**btrfs device** returns a zero exit status if it succeeds. Non zero is returned in case of failure.

If the *-c* option is used, *btrfs device stats* will add 64 to the exit status if any of the error counters is non-zero.

## AVAILABILITY

**btrfs** is part of btrfs-progs. Please refer to the documentation at <https://btrfs.readthedocs.io>.

## SEE ALSO

[btrfs-balance\(8\)](#) [btrfs-device\(8\)](#), [btrfs-replace\(8\)](#), [mkfs.btrfs\(8\)](#)

Previous

Next

© Copyright .

Built with [Sphinx](#) using a [theme](#) provided by [Read the Docs](#).