

btrfs-check(8)
btrfs-convert(8)
btrfs-device(8)
btrfs-filesystem(8)
btrfs-find-root(8)
btrfs-image(8)
btrfs-inspect-internal(8)

▾ btrfs-ioctl(2)

NAME

DESCRIPTION

DATA STRUCTURES AND DEFINITIONS

OVERVIEW

LIST OF IOCTLs

DETAILED DESCRIPTION

AVAILABILITY

SEE ALSO

btrfs-map-logical(8)
btrfs-property(8)
btrfs-qgroup(8)
btrfs-quota(8)
btrfs-receive(8)
btrfs-replace(8)
btrfs-rescue(8)
btrfs-restore(8)
btrfs-scrub(8)
btrfs-select-super(8)
btrfs-send(8)
btrfs-subvolume(8)
btrfstune(8)
fsck.btrfs(8)
mkfs.btrfs(8)

Administration
Hardware considerations
Changes (feature/version)
Changes (kernel/version)
Changes (btrfs-progs)
Contributors
Glossary
Installation instructions
Source repositories
Interoperability

FEATURES

Common Linux features
Custom ioctl's
Auto-repair on read
Balance
Compression
Checksumming
Convert
Deduplication
Defragmentation
Inline files
Quota groups
Reflink
Resize
Scrub
Seeding device
Send/receive
Subpage support
Subvolumes
Swapfile
Tree checker
Trim/discard
Volume management
Zoned mode

DEVELOPER DOCUMENTATION

Development notes
Developer's FAQ
Conventions and style for documentation
Experimental features
Btrfs design
Btrees
On-disk Format
Send stream format
JSON output
Internal APIs
Release checklist
Pull request review workflow
Command line, formatting, UI guidelines

▾ btrfs-ioctl(2)

NAME

DESCRIPTION

DATA STRUCTURES AND DEFINITIONS

OVERVIEW

LIST OF IOCTLs

DETAILED DESCRIPTION

AVAILABILITY

SEE ALSO

TODO
Troubleshooting pages

btrfs-ioctl(2)

NAME

btrfs-ioctl - documentation for the ioctl interface to btrfs

DESCRIPTION

The `ioctl()` system call is a way how to request custom actions performed on a filesystem beyond the standard interfaces (like `syscalls`). An `ioctl` is specified by a number and an associated data structure that implement a feature, usually not available in other filesystems. The number of `ioctl`s grows over time and in some cases get promoted to a VFS-level `ioctl` once other filesystems adopt the functionality. Backward compatibility is maintained and a formerly private `ioctl` number could become available on the VFS level.

DATA STRUCTURES AND DEFINITIONS

```
struct btrfs_ioctl_vol_args {
    __s64 fd;
    char name[BTRFS_PATH_NAME_MAX + 1];
};

#define BTRFS_SUBVOL_RDONLY          (1ULL << 1)
#define BTRFS_SUBVOL_QGROUP_INHERIT (1ULL << 2)
#define BTRFS_DEVICE_SPEC_BY_ID     (1ULL << 3)
#define BTRFS_SUBVOL_SPEC_BY_ID     (1ULL << 4)

struct btrfs_ioctl_vol_args_v2 {
    __s64 fd;
    __u64 transid;
    __u64 flags;
    union {
        struct {
            __u64 size;
            struct btrfs_qgroup_inherit __user *qgroup_inherit;
        };
        __u64 unused[4];
    };
    union {
        char name[BTRFS_SUBVOL_NAME_MAX + 1];
        __u64 devid;
        __u64 subvolid;
    };
};

struct btrfs_ioctl_get_subvol_info_args {
    /* Id of this subvolume */
    __u64 treeid;

    /* Name of this subvolume, used to get the real name at mount point */
    char name[BTRFS_VOL_NAME_MAX + 1];

    /*
     * Id of the subvolume which contains this subvolume.
     * Zero for top-level subvolume or a deleted subvolume.
     */
    __u64 parent_id;

    /*
     * Inode number of the directory which contains this subvolume.
     * Zero for top-level subvolume or a deleted subvolume
     */
    __u64 dirid;

    /* Latest transaction id of this subvolume */
    __u64 generation;

    /* Flags of this subvolume */
    __u64 flags;

    /* UUID of this subvolume */
    __u8 uuid[BTRFS_UUID_SIZE];

    /*
     * UUID of the subvolume of which this subvolume is a snapshot.
     * ALL zero for a non-snapshot subvolume.
     */
    __u8 parent_uuid[BTRFS_UUID_SIZE];

    /*
     * UUID of the subvolume from which this subvolume was received.
     * ALL zero for non-received subvolume.
     */
    __u8 received_uuid[BTRFS_UUID_SIZE];

    /* Transaction id indicating when change/create/send/receive happened */
    __u64 ctransid;
    __u64 otransid;
    __u64 stransid;
    __u64 rtransid;
    /* Time corresponding to c/o/s/rtransid */
    struct btrfs_ioctl_timespec ctime;
    struct btrfs_ioctl_timespec otime;
    struct btrfs_ioctl_timespec stime;
    struct btrfs_ioctl_timespec rtime;

    /* Must be zero */
    __u64 reserved[8];
};

#define BTRFS_QGROUP_INHERIT_SET_LIMITS (1ULL << 0)

struct btrfs_qgroup_inherit {
    __u64 flags;
    __u64 num_qgroups;
    __u64 num_ref_copies;
    __u64 num_excl_copies;
    struct btrfs_qgroup_limit lim;
    __u64 qgroups[];
};

#define BTRFS_QGROUP_LIMIT_MAX_RFER (1ULL << 0)
#define BTRFS_QGROUP_LIMIT_MAX_EXCL (1ULL << 1)
#define BTRFS_QGROUP_LIMIT_RSV_RFER (1ULL << 2)
#define BTRFS_QGROUP_LIMIT_RSV_EXCL (1ULL << 3)
#define BTRFS_QGROUP_LIMIT_RFER_CMPR (1ULL << 4)
#define BTRFS_QGROUP_LIMIT_EXCL_CMPR (1ULL << 5)

struct btrfs_qgroup_limit {
    __u64 flags;
    __u64 max_rfer;
    __u64 max_excl;
    __u64 rsv_rfer;
    __u64 rsv_excl;
};

/* Request information about checksum type and size */
#define BTRFS_FS_INFO_FLAG_CSUM_INFO (1U << 0)
/* Request information about filesystem generation */
#define BTRFS_FS_INFO_FLAG_GENERATION (1U << 1)
/* Request information about filesystem metadata UUID */
#define BTRFS_FS_INFO_FLAG_METADATA_UUID (1U << 2)

struct btrfs_ioctl_fs_info_args {
    __u64 max_id; /* out */
    __u64 num_devices; /* out */
    __u8 fsid[BTRFS_FSID_SIZE]; /* out */
    __u32 nodesize; /* out */
    __u32 sectorsize; /* out */
    __u32 clone_alignment; /* out */
    /* See BTRFS_FS_INFO_FLAG_* */
    __u16 csum_type; /* out */
    __u16 csum_size; /* out */
    __u64 flags; /* in/out */
    __u64 generation; /* out */
    __u8 metadata_uuid[BTRFS_FSID_SIZE]; /* out */
    __u8 reserved[944]; /* pad to 1k */
};
```

btrfs-check(8)
btrfs-convert(8)
btrfs-device(8)
btrfs-filesystem(8)
btrfs-find-root(8)
btrfs-image(8)
btrfs-inspect-internal(8)
btrfs-ioctl(2)
NAME
DESCRIPTION
DATA STRUCTURES AND DEFINITIONS
OVERVIEW
LIST OF IOCTLS
DETAILED DESCRIPTION
AVAILABILITY
SEE ALSO
btrfs-map-logical(8)
btrfs-property(8)
btrfs-qgroup(8)
btrfs-quota(8)
btrfs-receive(8)
btrfs-replace(8)
btrfs-rescue(8)
btrfs-restore(8)
btrfs-scrub(8)
btrfs-select-super(8)
btrfs-send(8)
btrfs-subvolume(8)
btrfstune(8)
fsck.btrfs(8)
mkfs.btrfs(8)

Administration
Hardware considerations
Changes (feature/version)
Changes (kernel/version)
Changes (btrfs-progs)
Contributors
Glossary
Installation instructions
Source repositories
Interoperability

FEATURES
Common Linux features
Custom ioctls
Auto-repair on read
Balance
Compression
Checksumming
Convert
Deduplication
Defragmentation
Inline files
Quota groups
Reflink
Resize
Scrub
Seeding device
Send/receive
Subpage support
Subvolumes
Swapfile
Tree checker
Trim/discard
Volume management
Zoned mode

DEVELOPER DOCUMENTATION
Development notes
Developer's FAQ
Conventions and style for documentation
Experimental features
Btrfs design
Btrees
On-disk Format
Send stream format
JSON output
Internal APIs
Release checklist
Pull request review workflow
Command line, formatting, UI guidelines

btrfs-ioctl(2)
NAME
DESCRIPTION
DATA STRUCTURES AND DEFINITIONS
OVERVIEW
LIST OF IOCTLS
DETAILED DESCRIPTION
AVAILABILITY
SEE ALSO

TODO
Troubleshooting pages

<pre>#define BTRFS_INO_LOOKUP_PATH_MAX 4080 struct btrfs_ioctl_ino_lookup_args { __u64 treeid; __u64 objectid; char name[BTRFS_INO_LOOKUP_PATH_MAX]; };</pre>
--

<pre>/* Specify the subvolid. */ #define BTRFS_SUBVOL_SYNC_WAIT_FOR_ONE (0) /* Wait for all currently queued. */ #define BTRFS_SUBVOL_SYNC_WAIT_FOR_QUEUED (1) /* Count number of queued subvolumes. */ #define BTRFS_SUBVOL_SYNC_COUNT (2) /* * Read which is the first in the queue (to be cleaned or being cleaned already), * or 0 if the queue is empty. */ #define BTRFS_SUBVOL_SYNC_PEEK_FIRST (3) /* Read the last subvolid in the queue, or 0 if the queue is empty. */ #define BTRFS_SUBVOL_SYNC_PEEK_LAST (4) struct btrfs_ioctl_subvol_wait { __u64 subvolid; __u32 mode; __u32 count; };</pre>

Constant name	Value
BTRFS_UUID_SIZE	16
BTRFS_FSID_SIZE	16
BTRFS_SUBVOL_NAME_MAX	4039
BTRFS_PATH_NAME_MAX	4087
BTRFS_VOL_NAME_MAX	255
BTRFS_LABEL_SIZE	256
BTRFS_FIRST_FREE_OBJECTID	256

OVERVIEW

The ioctls are defined by a number and associated with a data structure that contains further information. All ioctls use file descriptor (fd) as a reference point, it could be the filesystem or a directory inside the filesystem.

An ioctl can be used in the following schematic way:

<pre>struct btrfs_ioctl_args args; memset(&args, 0, sizeof(args)); args.key = value; ret = ioctl(fd, BTRFS_IOC_NUMBER, &args);</pre>

The 'fd' is the entry point to the filesystem and for most ioctls it does not matter which file or directory is that. Where it matters it's explicitly mentioned. The 'args' is the associated data structure for the request. It's strongly recommended to initialize the whole structure to zeros as this is future-proof when the ioctl gets further extensions. Not doing that could lead to mismatch of old userspace and new kernel versions, or vice versa. The 'BTRFS_IOC_NUMBER' is says which operation should be done on the given arguments. Some ioctls take a specific data structure, some of them share a common one, no argument structure ioctls exist too.

The library *libbtrfsutil* wraps a few ioctls for convenience. Using raw ioctls is not discouraged but may be cumbersome though it does not need additional library dependency. Backward compatibility is guaranteed and incompatible changes usually lead to a new version of the ioctl. Enhancements of existing ioctls can happen and depend on additional flags to be set. Zeroed unused space is commonly understood as a mechanism to communicate the compatibility between kernel and userspace and thus *zeroing is really important*. In exceptional cases this is not enough and further flags need to be passed to distinguish between zero as implicit unused initialization and a valid zero value. Such cases are documented.

File descriptors of regular files are obtained by `int fd = open()`, directories opened as `DIR *dir = opendir()` can be converted to the corresponding file descriptor by `fd = dirfd(dir)`.

LIST OF IOCTLS

Name	Description	Data
BTRFS_IOC_SNAP_CREATE	(obsolete) create a snapshot of a subvolume	struct btrfs_ioctl_vol_args
BTRFS_IOC_DEFRAG		
BTRFS_IOC_RESIZE		
BTRFS_IOC_SCAN_DEV	scan and register a given device path with filesystem module	struct btrfs_ioctl_vol_args
BTRFS_IOC_SYNC	Sync the filesystem, possibly process queued up work	NULL
BTRFS_IOC_CLONE		
BTRFS_IOC_ADD_DEV	add a device to the filesystem by path	struct btrfs_ioctl_vol_args
BTRFS_IOC_RM_DEV	delete a device from the filesystem by path	struct btrfs_ioctl_vol_args
BTRFS_IOC_BALANCE		
BTRFS_IOC_CLONE_RANGE		
BTRFS_IOC_SUBVOL_CREATE	(obsolete) create a subvolume	struct btrfs_ioctl_vol_args
BTRFS_IOC_SNAP_DESTROY	(obsolete) delete a subvolume	struct btrfs_ioctl_vol_args
BTRFS_IOC_DEFRAG_RANGE		
BTRFS_IOC_TREE_SEARCH		
BTRFS_IOC_TREE_SEARCH_V2		
BTRFS_IOC_INO_LOOKUP	resolve inode number to path, or lookup containing subvolume id	struct btrfs_ioctl_ino_lookup_args
BTRFS_IOC_DEFAULT_SUBVOL	set the default subvolume id	uint64_t
BTRFS_IOC_SPACE_INFO		
BTRFS_IOC_START_SYNC		
BTRFS_IOC_WAIT_SYNC		
BTRFS_IOC_SNAP_CREATE_V2	create a snapshot of a subvolume	struct btrfs_ioctl_vol_args_v2
BTRFS_IOC_SUBVOL_CREATE_V2	create a subvolume	struct btrfs_ioctl_vol_args_v2
BTRFS_IOC_SUBVOL_GETFLAGS	get flags of a subvolume	uint64_t
BTRFS_IOC_SUBVOL_SETFLAGS	set flags of a subvolume	uint64_t
BTRFS_IOC_SCRUB		
BTRFS_IOC_SCRUB_CANCEL		
BTRFS_IOC_SCRUB_PROGRESS		
BTRFS_IOC_DEV_INFO		
BTRFS_IOC_FS_INFO	get information about filesystem (device count, fsid, ...)	struct btrfs_ioctl_fs_info_args
BTRFS_IOC_BALANCE_V2		
BTRFS_IOC_BALANCE_CTL		
BTRFS_IOC_BALANCE_PROGRESS		
BTRFS_IOC_INO_PATHS		
BTRFS_IOC_LOGICAL_INO		
BTRFS_IOC_SET_RECEIVED_SUBVOL		
BTRFS_IOC_SEND		
BTRFS_IOC_DEVICES_READY		

btrfs-check(8)

btrfs-convert(8)

btrfs-device(8)

btrfs-filesystem(8)

btrfs-find-root(8)

btrfs-image(8)

btrfs-inspect-internal(8)

btrfs-ioctl(2)

NAME

DESCRIPTION

DATA STRUCTURES AND DEFINITIONS

OVERVIEW

LIST OF IOCTLs

DETAILED DESCRIPTION

AVAILABILITY

SEE ALSO

btrfs-map-logical(8)

btrfs-property(8)

btrfs-qgroup(8)

btrfs-quota(8)

btrfs-receive(8)

btrfs-replace(8)

btrfs-rescue(8)

btrfs-restore(8)

btrfs-scrub(8)

btrfs-select-super(8)

btrfs-send(8)

btrfs-subvolume(8)

btrfstune(8)

fsck.btrfs(8)

mkfs.btrfs(8)

Administration

Hardware considerations

Changes (feature/version)

Changes (kernel/version)

Changes (btrfs-progs)

Contributors

Glossary

Installation instructions

Source repositories

Interoperability

FEATURES

Common Linux features

Custom ioctl's

Auto-repair on read

Balance

Compression

Checksumming

Convert

Deduplication

Defragmentation

Inline files

Quota groups

Reflink

Resize

Scrub

Seeding device

Send/receive

Subpage support

Subvolumes

Swapfile

Tree checker

Trim/discard

Volume management

Zoned mode

DEVELOPER DOCUMENTATION

Development notes

Developer's FAQ

Conventions and style for documentation

Experimental features

Btrfs design

Btrees

On-disk Format

Send stream format

JSON output

Internal APIs

Release checklist

Pull request review workflow

Command line, formatting, UI guidelines

btrfs-ioctl(2)

NAME

DESCRIPTION

DATA STRUCTURES AND DEFINITIONS

OVERVIEW

LIST OF IOCTLs

DETAILED DESCRIPTION

AVAILABILITY

SEE ALSO

TODO

Troubleshooting pages

Field	Description
ioctl fd	file descriptor of the parent directory of the new subvolume
ioctl args	struct btrfs_ioctl_vol_args
args.fd	ignored
args.name	name of the subvolume, although the buffer can be almost 4KiB, the file size is limited by Linux VFS to 255 characters and must not contain a slash ('/')

BTRFS_IOC_SNAP_DESTROY

Note

obsoleted by [BTRFS_IOC_SNAP_DESTROY_V2](#)

(since: 2.6.33, obsoleted: 5.7) Delete a subvolume.

Field	Description
ioctl fd	file descriptor of the parent directory of the new subvolume
ioctl args	struct btrfs_ioctl_vol_args
args.fd	ignored
args.name	name of the subvolume, although the buffer can be almost 4KiB, the file size is limited by Linux VFS to 255 characters and must not contain a slash ('/')

BTRFS_IOC_INO_LOOKUP

Resolve inode number to a path (requires CAP_SYS_ADMIN), or read a containing subvolume id of the given file (unrestricted, special case). The size of the path name buffer is shorter than *PATH_MAX* (4096), it's possible that the path is trimmed due to that. Also implemented by [btrfs inspect-internal rootid](#).

The general case needs CAP_SYS_ADMIN and can resolve any file to its path. The special case for reading the containing subvolume is not restricted:

```
struct btrfs_ioctl_ino_lookup_args args;

fd = open("file", ...);
args.treeid = 0;
args.objectid = BTRFS_FIRST_FREE_OBJECTID;
ioctl(fd, BTRFS_IOC_INO_LOOKUP, &args);
/* args.treeid now contains the subvolume id */
```

Field	Description
ioctl fd	file descriptor of the file or directory to lookup the subvolumeid
ioctl args	struct btrfs_ioctl_ino_lookup_args
args.treeid	subvolume id against which the path should be resolved (needs CAP_SYS_ADMIN), or 0 so the subvolume containing <i>fd</i> will be used
args.objectid	inode number to lookup, <i>INODE_REF_KEY</i> with that key.objectid, or BTRFS_FIRST_FREE_OBJECTID as special case to read only the tree id and clear the <i>args.name</i> buffer
args.name	path relative to the toplevel subvolume, or empty string

BTRFS_IOC_DEFAULT_SUBVOL

Set the given subvolume id as the default one when mounting the filesystem without *subvol=path* or *subvolid=id* options.

Field	Description
ioctl fd	file descriptor of the directory inside which to create the new snapshot
ioctl args	numeric value of subvolume to become default (uint64_t)

BTRFS_IOC_SNAP_CREATE_V2

Create a snapshot of a subvolume.

Field	Description
ioctl fd	file descriptor of the directory inside which to create the new snapshot
ioctl args	struct btrfs_ioctl_vol_args_v2
args.fd	file descriptor of any directory inside the subvolume to snapshot, must be on the filesystem
args.transid	ignored
args.flags	any subset of <i>BTRFS_SUBVOL_RDONLY</i> to make the new snapshot read-only, or <i>BTRFS_SUBVOL_QGROUP_INHERIT</i> to apply the <i>qgroup_inherit</i> field
args.name	the name, under the ioctl fd, for the new subvolume

BTRFS_IOC_SUBVOL_CREATE_V2

(since: 3.6) Create a subvolume, qgroup inheritance and limits can be specified.

Field	Description
ioctl fd	file descriptor of the parent directory of the new subvolume
ioctl args	struct btrfs_ioctl_vol_args_v2
args.fd	ignored
args.transid	ignored
args.flags	flags to set on the subvolume, <code>BTRFS_SUBVOL_RDONLY</code> for readonly, <code>BTRFS_SUBVOL_QGROUP_INHERIT</code> if the qgroup related fields should be processed
args.size	number of entries in <code>args.qgroup_inherit</code>
args.qgroup_inherit	inherit the given qgroups (struct btrfs_qgroup_inherit) and limits (struct btrfs_qgroup_limit)
name	name of the subvolume, although the buffer can be almost 4KiB, the file size is limited by Linux VFS to 255 characters and must not contain a slash ('/')

BTRFS_IOC_SUBVOL_GETFLAGS

Read the flags of a subvolume. The returned flags are either 0 or *BTRFS_SUBVOL_RDONLY*.

Field	Description
ioctl fd	file descriptor of the subvolume to examine
ioctl args	uint64_t

BTRFS_IOC_SUBVOL_SETFLAGS

Change the flags of a subvolume.

Field	Description
ioctl fd	file descriptor of the subvolume to modify
ioctl args	uint64_t, either 0 or <i>BTRFS_SUBVOL_RDONLY</i>

BTRFS_IOC_GET_FSLABEL

Read the label of the filesystem into a given buffer. Alternatively it can be read from `/sys/fs/btrfs/FSID/label` though it requires to know the FSID of the filesystem.

Field	Description
ioctl fd	file descriptor of any file/directory in the filesystem

[< Previous](#)

[Next >](#)

AVAILABILITY

btrfs is part of **btrfs-progs**. Please refer to the documentation at <https://btrfs.readthedocs.io>.

ioctl(2)